



Lab Guide: Sentiment Analysis and Brand Monitoring

- [Lab 1: Overview](#)
- [Lab 2: Twitter input](#)
- [Lab 3: External API and REST filter](#)
- [Lab 4: Elasticsearch Machine Learning](#)

Lab 1: Overview

Objective: In this lab, we are going to first create a twitter account and then prepare our environment for the rest of the class.

1. The first step will be to create a Twitter app. In order to do that you will need a twitter account. Use your twitter account or create a new one: <https://apps.twitter.com> to create an application.

The screenshot shows the Twitter Apps landing page. At the top right, there is a "Sign in" link. Below it, a message states: "By using Twitter's services you agree to our [Cookie Use](#) and [Data Transfer](#) outside the EU. We and our partners operate globally and use cookies, including for analytics, personalisation, and ads." A small "X" icon is at the end of this message. The main heading is "Twitter Apps", with a sub-instruction: "Please [sign in](#) with your Twitter Account to create and maintain Twitter Apps." At the bottom right of the main content area is a "Tweet" button. At the very bottom of the page, there are links for "About", "Terms", "Privacy", and "Cookies", along with the copyright notice: "© 2018 Twitter, Inc."

This screenshot is identical to the one above, showing the Twitter Apps landing page. The "Create New App" button is highlighted with a red circle.

This screenshot shows the "Application Details" form. Several fields are highlighted with red circles: the "Name" field containing "Elastic training", the "Description" field containing "Collecting tweet", the "Website" field containing "http://placeholder.com", and the "Developer Agreement" checkbox which is checked and highlighted. At the bottom of the form, there is a "Create your Twitter application" button.

Elastic training

Application Settings

Consumer Key (API Key) [REDACTED]

Consumer Secret (API Secret) [REDACTED]

Access Level Read and write (modify app permissions)

Owner JohnCen89795791

Owner ID 1013833651258306560

Application Actions

Regenerate Consumer Key and Secret Change App Permissions

Your Access Token

You haven't authorized this application for your own account yet.

By creating your access token here, you will have everything you need to make API calls right away. The access token generated will be assigned your application's current permission level.

Token Actions

Create my access token

2. Let's start by going into server1:

```
ssh server1
```

3. Extract the two folders:

```
tar -zxf elasticsearch-6.3.1.tar.gz
tar -zxf kibana-6.3.1-linux-x86_64.tar.gz
```

4. You can now start elasticsearch.

```
./elasticsearch-6.3.1/bin/elasticsearch -E network.host=0.0
```

5. We should as well extract Logstash.

```
ssh server2
```

Extract logstash:

```
tar -zxf logstash-6.3.1.tar.gz
```

6. In order to make sure that we can use our Twitter tokens without them being visible let's use the Logstash keystore:

```
./logstash-6.3.1/bin/logstash-keystore create
```

You can now add each one of your tokens. Make sure to not paste your token before the prompt tell you so in order to keep your tokens as secret as possible:

```
./logstash-6.3.1/bin/logstash-keystore add CONS_KEY
```

Then copy your consumer key. Reiterate the process for CONS_KEY_SEC (consumer key secret), OAUTH_TO (Access token), OAUTH_TO_SEC (Access token secret)

List all your tokens to make sure you didn't forgot one:

```
./logstash-6.3.1/bin/logstash-keystore list
```

7. We are going to install some Logstash plugins during this class. Since we don't have any SSL certificate we need to remove "https" in the Gemfile.

```
vim logstash-6.3.1/Gemfile
```

Replace "<https://rubygems.org>" by "<http://rubygems.org>"

8. You can now start Kibana on the server1.

```
./kibana-6.3.1-linux-x86_64/bin/kibana --host=0.0.0.0
```

9. We need to add a dataset that we are going to use in this training. In the server1 there is folder, datasets, with a subfolder called snapshot. We are going to use it in order to load our data. Open the **elasticsearch.yml** file (in folder **.elasticsearch-6.3.1/config**) and add the following line:

```
path.repo: "/home/elasticsearch/datasets/snapshot"
```

Restart your elasticsearch node. To stop it, go into the terminal tab where it is running and press **ctrl + c**, then restart it.

10. Once Elasticsearch is running go into Kibana go into **Dev Tools** and type the following:

```
PUT /_snapshot/tweet
{
  "type": "fs",
  "settings": {
    "location": "/home/elasticsearch/datasets/snapshot"
  }
}
```

This is defining the repository for the snapshot. Now let's use the snapshot to restore the dataset:

```
POST /_snapshot/tweet/_dump/_restore
```

11. In the Kibana console verify if the snapshot worked properly:

```
GET _cat/indices
```

You should see an index called twitter-ml-clean.

12. We are going to use Elasticsearch Machine Learning in this class. To do so we need to activate the Elasticsearch license.

The screenshot shows the Kibana Management interface. On the left is a sidebar with icons for Discover, Visualize, Dashboard, Timeline, Canvas, APM, Dev Tools, Monitoring, and Management. The Management icon is highlighted with a blue background. The main area has a header "Management" and "Version: 6.3.1". Below this is a section for "Elasticsearch" with two tabs: "Index Management" and "License Management", with "License Management" circled in red. At the bottom are tabs for "Index Patterns", "Saved Objects", "Reporting", and "Advanced Settings".

The screenshot shows the "Management / Elasticsearch / License Management" page. The sidebar is identical to the previous one. The main content area starts with a message: "Your Basic license is active" with a note "Your license will never expire." Below this are two boxes: "Update your license" (with a "Update license" button) and "Start a 30-day trial" (with a "Start trial" button circled in red). The "Start trial" button is located in a box that also contains the text "Experience what security, machine learning, and all our other Platinum features have to offer."

Summary: In this lab you should have created a twitter application, which should give you access to different tokens and you should have activated the free license and set up your cluster for the next labs!

End of Lab 1

Lab 2: Twitter input

Objective: In this lab, you are going to use the Twitter account you just created in order to collects tweets using different methodologies.

1. In the folder **pipelines** you are going to find a file called **lab2.conf**, edit this file. But before that you will need to ssh into server2.

```
ssh server2
vim pipelines/lab2.conf
```

Let's change the input in order to collect random tweets.

```
input {
  twitter {
    consumer_key => "${CONS_KEY}"
    consumer_secret => "${CONS_KEY_SEC}"
    oauth_token => "${OAUTH_TO}"
    oauth_token_secret => "${OAUTH_TO_SEC}"
    use_samples => true
    full_tweet => true
    type => "tweet" # Apply a new field to the JSON collection
    # it is pretty useful when working with multiple inputs
  }
}
```

2. You probably realised that we used `prune` in the pipeline, so we need to install it. Hopefully a single line will do the trick:

```
logstash-6.3.1/bin/logstash-plugin install logstash-filter-prune
```

It should take only few seconds. Once it's done, launch the pipeline.

```
./logstash-6.3.1/bin/logstash -f pipelines/lab2.conf
```

3. Let's see if our pipeline is collecting tweets properly. Wait a little bit and then Go in **Kibana** and count the number of documents present in the **twitter** index.

[Hide answer:](#)

```
GET twitter/_count
```

4. Let's stop the pipeline and collect tweets based on other criterion. Delete the **twitter** index.

[Hide answer:](#)

In the terminal where your pipeline is running press **ctrl + c**.

In the Kibana console:

```
DELETE twitter
```

5. Let's collect tweets from specific area. To do this we are going to use the location contained in a tweet. We are going to collect tweets from the US west coast, UK and Australia. For every box we need to define 2 points, The SW one and the NE one.

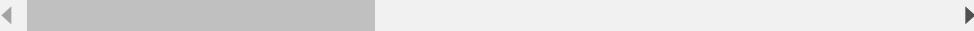
- For the US west coast:
 - SW: lon: -124.264886 lat: 31.598581
 - NE: lon:-117.570610 lat: 48.536975
- For the UK:
 - SW: lon: -12.006055 lat: 49.984609
 - NE: lon: 1.220102 lat: 59.406623
- For Australia:
 - SW: lon: 112.310377 lat: -39.696116

- NE: lon: 154.248758 lat: -10.751004

Update the configuration file to collect tweets from those area.

[Hide answer:](#)

```
input {  
    twitter {  
        consumer_key => "${CONS_KEY}"  
        consumer_secret => "${CONS_KEY_SEC}"  
        oauth_token => "${OAUTH_TO}"  
        oauth_token_secret => "${OAUTH_TO_SEC}"  
        locations => "-124.264886,31.598581,-117.570610,4  
        full_tweet => true  
        type => "tweet"  
    }  
}
```



6. Let's verify if the tweets that we are collecting really comes from the desired location. To do that we are going to build a coordinate map in **Kibana**. All the points should be on the US west coast, Australia or UK.

a. The first step will be to create an index pattern, go into management and then select index pattern:

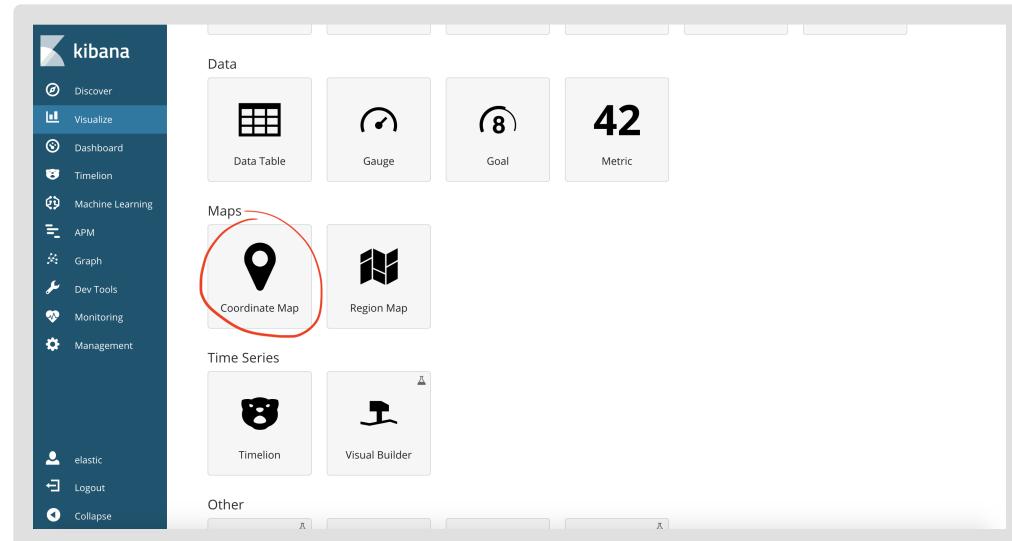
The screenshot shows the Kibana Management interface. On the left, a sidebar lists various features: Discover, Visualize, Dashboard, Timeline, Machine Learning, APM, Graph, Dev Tools, Monitoring, and Management. The 'Management' option is highlighted with a red circle. The main area is titled 'Management' and shows 'Version: 6.2.2'. It contains sections for Elasticsearch, License Management, Watcher, Kibana, Logstash, and Pipelines. The 'Index Patterns' tab is selected, also highlighted with a red circle.

- b. Type the name of the index pattern **twitter** and select the field **tweet_created_at** as the timefield:

- c. Click on **Visualize** and then the **+** button.

The screenshot shows the Kibana Visualize interface. On the left, the sidebar has 'Discover' and 'Visualize' selected, both highlighted with red circles. The main area is titled 'Visualize' and contains a search bar and a '+' button, which is also circled in red. Below these are several visualization options listed in pairs: Title (Data Table), author (Tag Cloud), New Visualization (Pie), retweet (Vertical Bar), tweet_histogram (Coordinate Map), and tweet_position (Coordinate Map). At the bottom, it says '0 items selected' and '1-5 of 5'.

- d. Choose **Coordinate map**



e. Select **twitter** index template

The Kibana Visualize interface shows the 'Choose search source' step. The sidebar still has 'Visualize' selected. The main area is split into two sections: 'From a New Search, Select Index' and 'Or, From a Saved Search'. In the 'From a New Search' section, there is a dropdown menu with three items: 'twitter-ml-original', 'twitter-ml', and 'twitter' (which is circled in red). In the 'Or, From a Saved Search' section, there is a dropdown menu with two items: 'negative' and 'positive'. There are also buttons for 'Manage saved searches' and '1-2 of 2'.

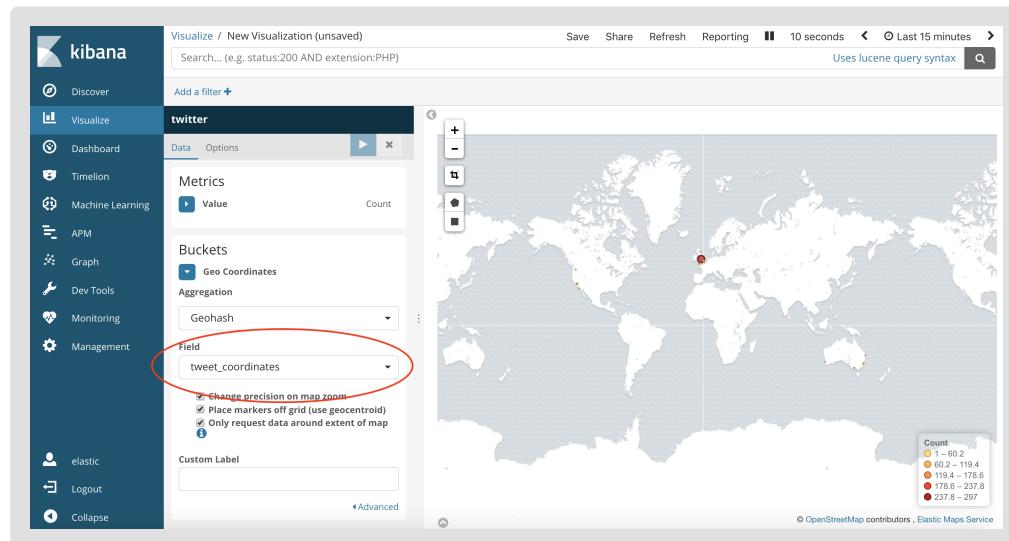
f. Click on **Geo Coordinates**

The screenshot shows the Kibana Visualize interface for a 'twitter' visualization. On the left, the sidebar includes options like Discover, Visualize (which is selected), Dashboard, Timelion, Machine Learning, APM, Graph, Dev Tools, Monitoring, and Management. The main area shows a world map with country boundaries. In the top right, there are buttons for Save, Share, Refresh, Reporting, and a time range selector set to 'Last 15 minutes'. Below these are buttons for 'Data' and 'Options'. Under the 'Metrics' section, there is a dropdown menu with 'Value' and 'Count' options. The 'Buckets' section has a dropdown menu labeled 'Select buckets type' with 'Geo Coordinates' highlighted and circled in red. A 'Cancel' button is also visible.

g. Select the aggregation Geohash

This screenshot continues from the previous one, showing the 'Aggregation' dropdown menu. The 'Geo Coordinates' option is still circled in red. The 'Select an aggregation' dropdown now has 'Geohash' highlighted and selected, indicated by a blue background. The rest of the interface remains the same, including the world map and the top navigation bar.

h. Select the field tweet_coordinates



i. Click on the play button

All the point should be on the US west coast, UK and Australia.

7. You can now stop the pipeline using **ctrl + c** in the terminal.

8. Now let's try to collect tweets from a specific user. Go on this website: <http://mytwitterid.com> and type your twitter pseudo. It will give you an ID. Use this ID to build a pipeline that will collect your tweets. You can find your id in the twitter app that you created earlier. Once you have saved the configuration file, start the pipeline.

[Hide answer:](#)

```
input {
  twitter {
    consumer_key => "${CONS_KEY}"
    consumer_secret => "${CONS_KEY_SEC}"
    oauth_token => "${OAUTH_TO}"
    oauth_token_secret => "${OAUTH_TO_SEC}"
    follows => ["YOUR_ID"]
    full_tweet => true
  }
}
```

```
        type => "tweet"
    }
}
```

9. Now go on twitter and tweet the following:

"Testing the twitter input for Logstash!"

<https://www.elastic.co/guide/en/logstash/current/plugins-inputs-twitter.html> #ElasticTraining #Elastic".

10. Now search for your tweet into your Elasticsearch using your ID and the field **author_id**.

```
GET twitter/_search
{
  "query": {
    "match": {
      "author_id": "YOUR_TWITTER_ID"
    }
  }
}
```

11. Let's try to collect tweets that contained at least one of the two following word: **HBO** or **Netflix**.

```
input {
  twitter {
    consumer_key => "${CONS_KEY}"
    consumer_secret => "${CONS_KEY_SEC}"
    oauth_token => "${OAUTH_TO}"
    oauth_token_secret => "${OAUTH_TO_SEC}"
    keywords => ["Netflix", "HBO"]
  }
}
```

```
full_tweet => true  
type => "tweet"  
}  
}
```

12. Start your new pipeline and check if your are collecting tweets by counting the number of document in your index.

[Hide answer:](#)

```
GET twitter/_count
```

Summary: Well done!! In this lab, you used Logstash to collect tweets using different kind of filter. You also learned how to use Kibana to check if your way of collecting tweets worked properly.

End of Lab 2

Lab 3: External API and REST filter

Objective: In this lab, you will see how to build the external classifier and then how to send events to the external API using the rest filter.

1. First let's see what are the data is that we are going to use to train our algorithm. We are going to use a dataset included in the

NLTK. The first step will be to build our virtual environment for python that will include the different library that we are going to use.

Go into the **sentiment_api** folder in the server2:

```
cd sentiment_api
```

Now we will need to build the virtual environement:

```
virtualenv -p python3 venv  
source venv/bin/activate
```

2. The first line will build a virtual environment to avoid us having to install the library globally. The second command line will activate the virtual environement. Now you can install the different libraries that we are going to use.

```
pip install -r requirements.txt
```

This will automatically download and install all the libraries include in the requirements file.

Now that we have all the needed libraries, let's start exploring the dataset.

```
python3
```

The above command will start a python shell. We need to import the library that contains the dataset.

```
from nltk.corpus import movie_reviews  
import nltk
```

```
nltk.download('movie_reviews')
nltk.download('punkt')
```

The corpus being loaded, let's start to look at it. Movie_reviews is an object that contains different attributes and methods. We can use them to display the different labels present in the dataset.

```
movie_reviews.categories()
```

This line of code is going to output the following: ["pos", "neg"], "pos" is standing for positive and "neg" is standing for negative. Now if you want to display all the reviews that are positive you can type the following. If I want to retrieve all the positive review I can type:

```
movie_reviews.fileids("pos")
```

This will display all the files that contain a positive review. Note that every file starts with "pos". Let's look at one of them:

```
file = movie_reviews.fileids("pos")[0]
movie_reviews.raw(file)
```

Now you can see what's the content of the first positive review. Let's count how many positive and negative review we have:

```
print("Number of positive review:", len(movie_reviews.fileids("pos")))
print("Number of negative review:", len(movie_reviews.fileids("neg")))
```

Ideally you would prefer a dataset that will reflect the real-world proportions. Let's now train an algorithm that will take this previous datasets in order to label unseen data. The Text blob

library can use a classifier from NLTK: **NaiveBayesAnalyzer**. By default this classifier is going to use the previous corpus.

```
from textblob import TextBlob  
from textblob.sentiments import NaiveBayesAnalyzer  
blob = TextBlob("This movie is awesome", analyzer=NaiveBayesAnalyzer)  
blob.sentiment
```

So if I type "This movie is awesome" we are going to compute the probability of this text to be positive or negative depending on the word from the previous dataset. The output is:

Sentiment(classification='pos', p_pos=0.6928102445189113, p_neg=0.30718975548108857)

It tells you that the probability of being positive is around 0.69 and the probability of being negative is 0.31, so the text will be labeled as positive, since the probability is higher.

You can see that every time you are using the commands:

```
blob = TextBlob("This movie is really bad", analyzer=NaiveBayesAnalyzer)  
blob.sentiment
```

It will take few seconds to be executed. It is because every time it will need to recompute the probability based on the dataset, so what you can do is use a Blobber factory, that will help dry your code (https://en.wikipedia.org/wiki/Don%27t_repeat_yourself).

```
from textblob import Blobber  
text_blob = Blobber(analyzer=NaiveBayesAnalyzer())  
text_blob("This is a test").sentiment
```

Now only the first execution will take time. The others will be much faster.

Maybe, you would prefer using another dataset in order to do the sentiment analysis. An option can be to use your own data. So you are going to build your own classifier:

```
from textblob.classifiers import NaiveBayesClassifier  
train = [('This movie is just... Wow', 'pos'),  
         ('I just wasted my time', 'neg')]  
classifier = NaiveBayesClassifier(train)  
classifier.classify("Wow this movie is incredible!")
```

Keep in mind that the training dataset is extremely important. Don't ever expect an algorithm with bad training data to give you good results. Garbage in, Garbage out.

```
exit()
```

Exit the python shell.

The code for the API has already been written down. Let's just start the API.

```
python sentiment_server.py
```

Let's have a look at **sentiment_server.py**. The method that is interesting is **predict_a_tweet**, it expects a JSON with a field **submit**. The content of that field will then be send to a function **clean_text** to be cleaned of token like **@username**. Then using the trained algorithm we compute the probability of the tweet to be positive. If this value is higher than **0.55** the label will be positive, if the probability is between **0.55 and 0.45** then we are not sure enough and the label will be neutral, lastly if the probability is lower than **0.45** then the label is negative. We are

sending back a JSON containing two fields: **sentiment** and **sentiment_score**, the first field is the label: "positive", "negative" or "neutral" the second field is the probability of being positive.

Our webserver is up and running. It is now time to set up Logstash in order to send the tweets to the API. The first step will be to install the rest filter.

```
logstash-6.3.1/bin/logstash-plugin install logstash-filter-...
```

3. In the file **lab3.conf**, try to find the rest filter in the configuration file and edit it accordingly. We need to specify that Logstash needs to send the field **tweet_content** to the URL: <http://localhost:5000/predict> using a **POST** request. And we want to put the result in a field called **rest**.

[Hide answer:](#)

```
rest {  
    request => {  
        url => "http://localhost:5000/predict"  
        method => "POST"  
        params => {  
            "submit" => "%{tweet_content}"  
        }  
        headers => { "Content-Type" => "application/json" }  
    }  
    target => 'rest'  
}
```

Summary: Well done!! Now every tweet is send to the API to get a label, and has been enriched in Logstash.

End of Lab 3

Lab 4: Elasticsearch Machine Learning

Objective: Objective: In this last lab, we are going to see how we can use the Elasticsearch build-in machine learning features in order to detect anomalies in the data. Those could be a certain amount of tweets that are not normal for this time of the day, or a user that is tweeting too much in comparison with other users.

1. Do a **match_all** query on the **twitter-ml-clean** index.

[Hide answer:](#)

```
GET twitter-ml-clean/_search
{
  "query": {
    "match_all": {}
  }
}
```

Or you could simply do:

```
GET twitter-ml-clean/_search
```

2. Let's run our first Machine Learning job. We want to know if at some point in time the number of tweet over time present some anomaly in its pattern. Go into **Machine Learning**, click on **Create new job**, select the **twitter-ml-clean** index pattern,

choose **Single metric**. You can now configure the job. We want to group our documents by bucket of **10 minutes** and do a **count** on those bucket. Click on the button **use full twitter-ml-clean data**. Give a **name** to your job and press the button **Create job**, and then view a job: **View Results**

[Hide answer:](#)

The screenshot shows the Kibana interface under the 'Job Management' tab. On the left sidebar, 'Machine Learning' is selected. A red circle highlights the '+ Create new job' button at the top right of the main content area. Below it, the message 'No jobs configured' is displayed. The top navigation bar includes links for 'Job Management', 'Anomaly Explorer', 'Single Metric Viewer', and 'Settings'. The bottom navigation bar has links for 'Discover', 'Visualize', 'Dashboard', 'Timeline', 'Machine Learning', 'APM', 'Graph', 'Dev Tools', 'Monitoring', and 'Management'. User information and logout/collapse buttons are also present.

The screenshot shows the 'Create New Job' configuration page. The left sidebar is identical to the previous screenshot. The main area is divided into two sections: 'From a New Search, Select Index' and 'Or, From a Saved Search'. In the 'From a New Search' section, a search bar is labeled 'Filter...' with 'Name' and 'twitter-ml-original' selected. A red circle highlights 'twitter-ml'. Below this, there are other index names: 'negative' and 'positive'. In the 'Or, From a Saved Search' section, a search bar is labeled 'Saved Searches Filter...' with '1-2 of 2' and 'Manage saved searches' buttons. The user has selected the 'negative' and 'positive' saved searches.

Kibana

- Discover
- Visualize
- Dashboard
- Timeline
- Machine Learning**
- APM
- Graph
- Dev Tools
- Monitoring
- Management

elastic **Logout** **Collapse**

Machine Learning / Job Management / Create New Job

Create a job from the index pattern twitter-ml

Use a wizard

Use one of the wizards to create a machine learning job to find anomalies in your data.

 **Single metric**
Detect anomalies in a single time series.

 **Multi metric**
Detect anomalies in multiple metrics by splitting a time series by a categorical field.

 **Population**
Detect activity that is unusual compared to the behavior of the population.

Advanced

Use the full range of options to create a job for more advanced use cases.

Learn more about your data

If you're not sure what type of job to create, first explore the fields and metrics in your data.

Kibana

- Discover
- Visualize
- Dashboard
- Timeline
- Machine Learning**
- APM
- Graph
- Dev Tools
- Monitoring
- Management

elastic **Logout** **Collapse**

Machine Learning / Job Management / Create New Job / Single Metric Job

New job from index pattern twitter-ml

Chart interval: 1h **Use full twitter-ml data**

Aggregation  Select an aggregation **Count**

Field  Select a field

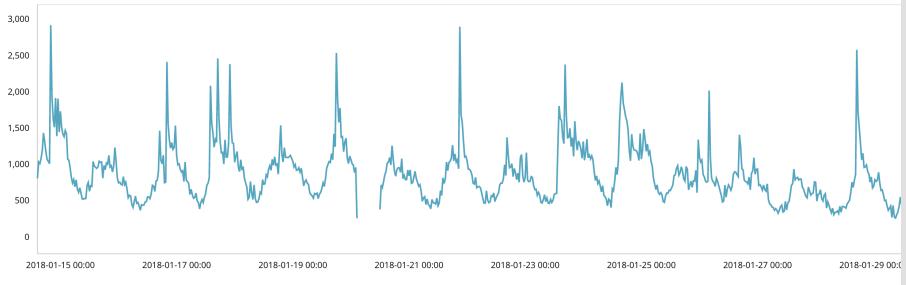
Bucket span  **10m** Estimate bucket span 

Kibana

- Discover
- Visualize
- Dashboard
- Timeline
- Machine Learning**
- APM
- Graph
- Dev Tools
- Monitoring
- Management

elastic **Logout** **Collapse**

Sparse data 



Name  **single_metric_job**

Description  Job description

Job Groups  Job Group

Advanced 

Create Job



3. Play with the time picker to navigate into your timeline. Why is the blue area at the end of the timeline different than at the beginning?

[Hide answer:](#)

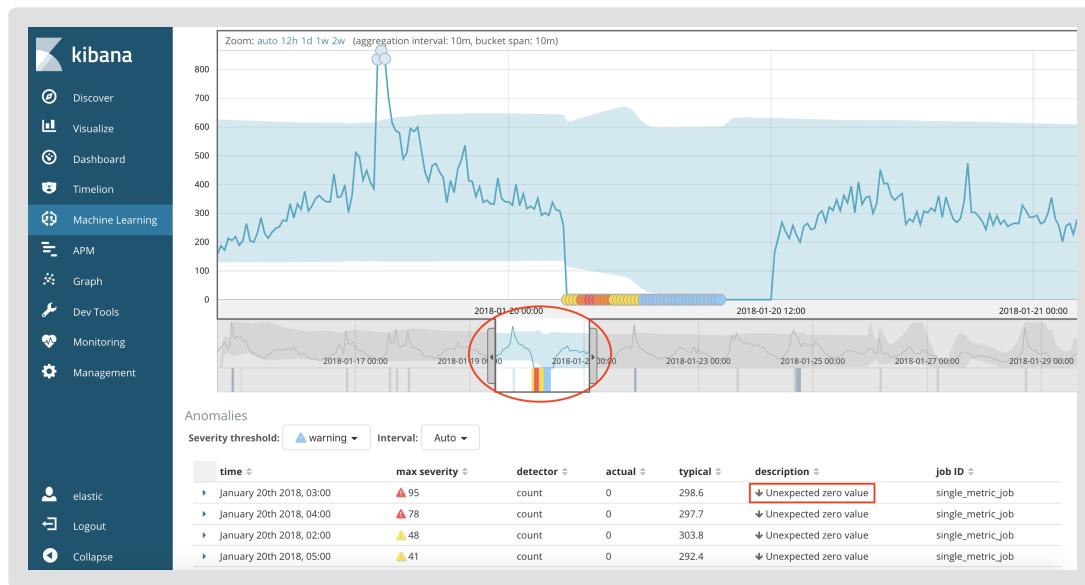
The time picker is surrounded by the red square. The blue area is the area in which your model expect the time series that you are analyzing to be in. At the beginning the model doesn't have enough data to define properly this area, at a certain point in time you have enough data to have an accurate model, so the blue area fit more accurately your time series.

4. Find the critical anomalies (The one in red). What are the description of those anomalies?

[Hide answer:](#)

Select the part of the time series that contains the critical anomalies and go into the anomalies listing. You can see on the right the reason of the anomalies. The anomaly should be "Unexpected zero value". Those anomalies are there because

the machine that was collecting data was down during this period of time.



5. Go to the last anomaly what is the "description" of this anomaly?

[Hide answer:](#)

This anomaly is due to values in the time serie 4 times higher than what it should "normally" be. That's means that a lot of tweets has been tweeted around this time.

6. We can identify when there is a change in the frequency of the events. Let's now leverage the sentiment that are present in our tweets. Redo the same steps than before but this time use the **multimetric job** instead. Select the field **event rate** and choose to build multiple time series based on the sentiment field. Group the documents by bucket of **10 minutes**, give a name to the job and then create the job.

[Hide answer:](#)

Create a job from the index pattern twitter-ml

Use a wizard
Use one of the wizards to create a machine learning job to find anomalies in your data.

- Single metric**
Detect anomalies in a single time series.
- Multi metric**
Detect anomalies in multiple metrics by splitting a time series by a categorical field. **(Red circle)**
- Advanced**
Use the full range of options to create a job for more advanced use cases.

Learn more about your data
If you're not sure what type of job to create, first explore the fields and metrics in your data.

Machine Learning / Job Management / Create New Job / Multi Metric job January 14th 2018, 14:33:44.000 to January 29th 2018, 17:01:44.000

New job from index pattern twitter-ml

Job settings

Fields

- event rate
- sentiment_score
- hashtags
- pseudo
- sentiment
- tweet_id
- Sparse data

Split Data: Remove split

Key Fields (Influencers)

Results

Document count

Data split by sentiment

positive

Count event rate

Machine Learning / Job Management / Create New Job / Multi Metric job January 14th 2018, 14:33:44.000 to January 29th 2018, 17:01:44.000

New job from index pattern twitter-ml

Job settings

Fields

- tweet_id
- Sparse data

Split Data: Remove split

Key Fields (Influencers)

Bucket span Estimate bucket span

Job Details

Name **(Red circle)**

Description

Job Groups

Advanced

Create Job **(Red circle)**

negative

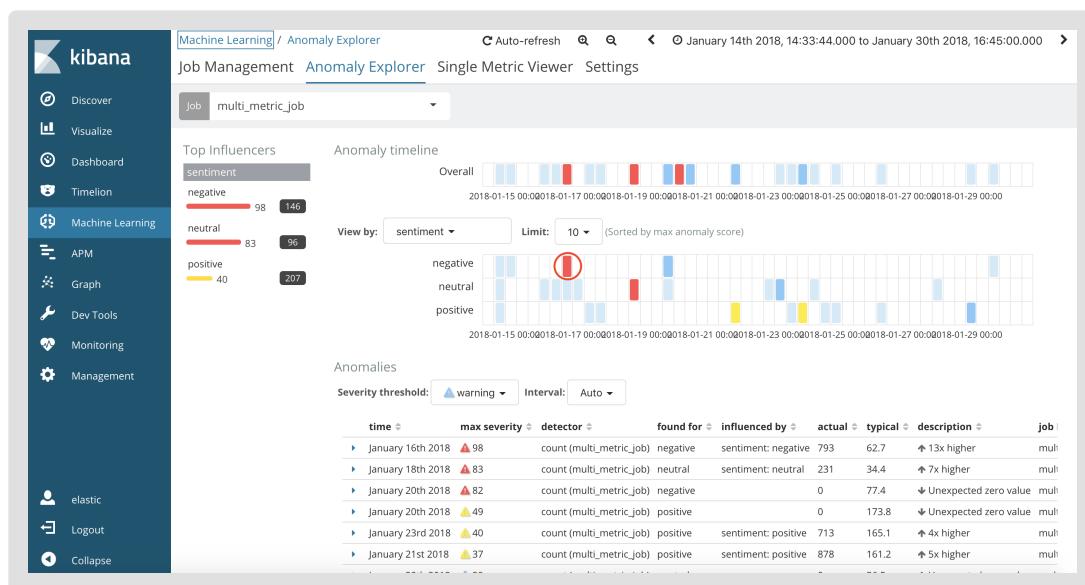
positive

Count event rate

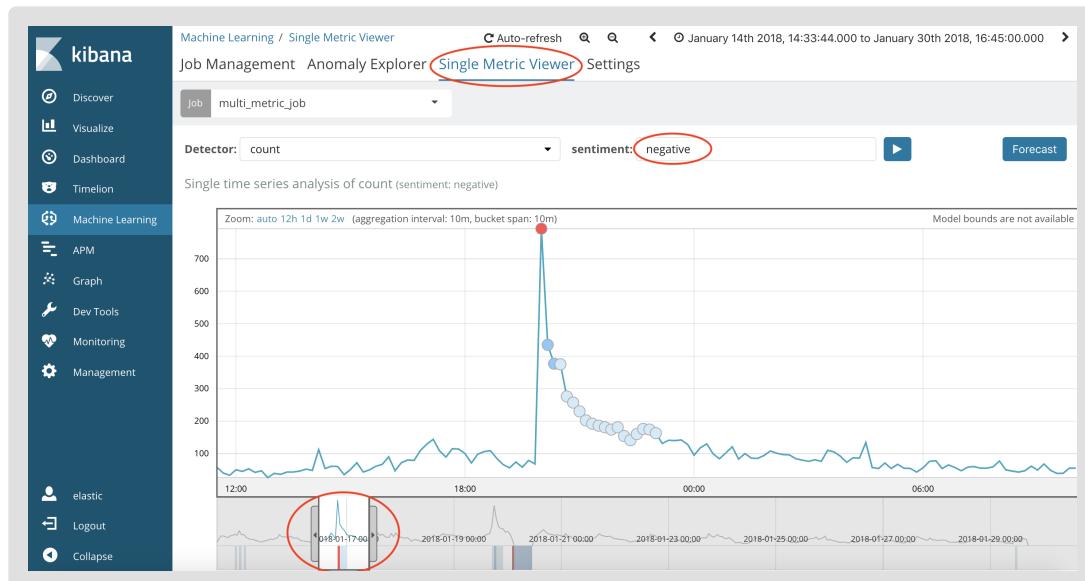
7. Click on **View Results**. The first page you see give you a summary of the different anomalies overall and a summary of the anomalies per time series. Let's say that you are interested in the negative tweets that you have been collecting. Find the first critical anomaly, from the "negative" time series?

[Hide answer:](#)

You can simply click on the anomaly:



Or you could click on the **Single metric viewer**, select the time series: **negative** and navigate directly to the anomaly.

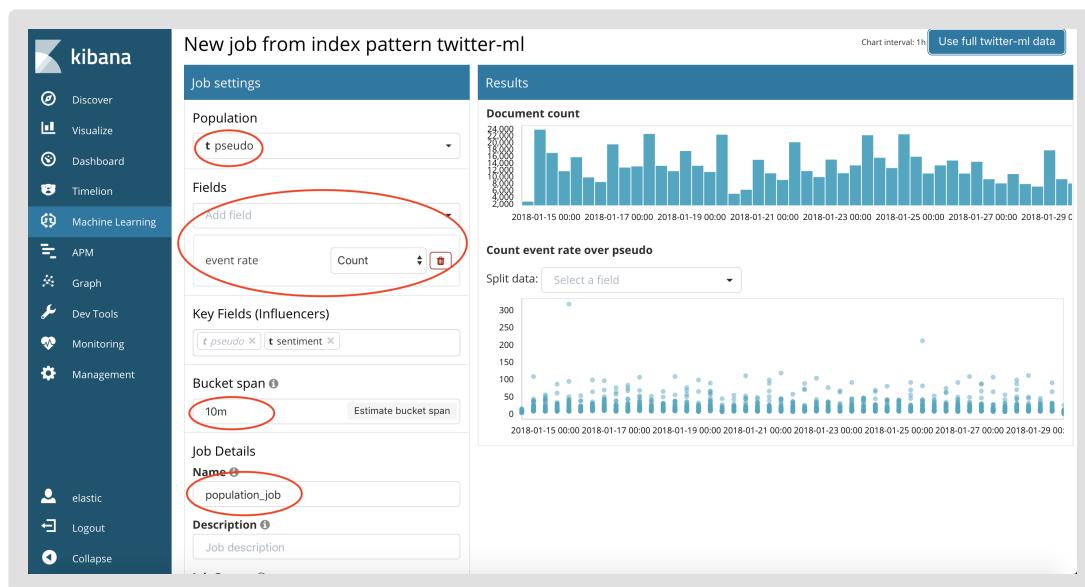


8. What is the **max severity**, the **actual** number of event in the bucket, the **typical** number of events and the description of the anomaly?

Show answer:

9. Let's now try to detect some users that have an abnormal behaviour in our data. To do so you are going to run a "Population" job, the entities that we want to monitor are the user's "pseudo". Use the **Population** job wizard, let's set the population to be **pseudo**, the field to be the **event rate**, a bucket of **10 minutes** and give a name to the job.

Hide answer:



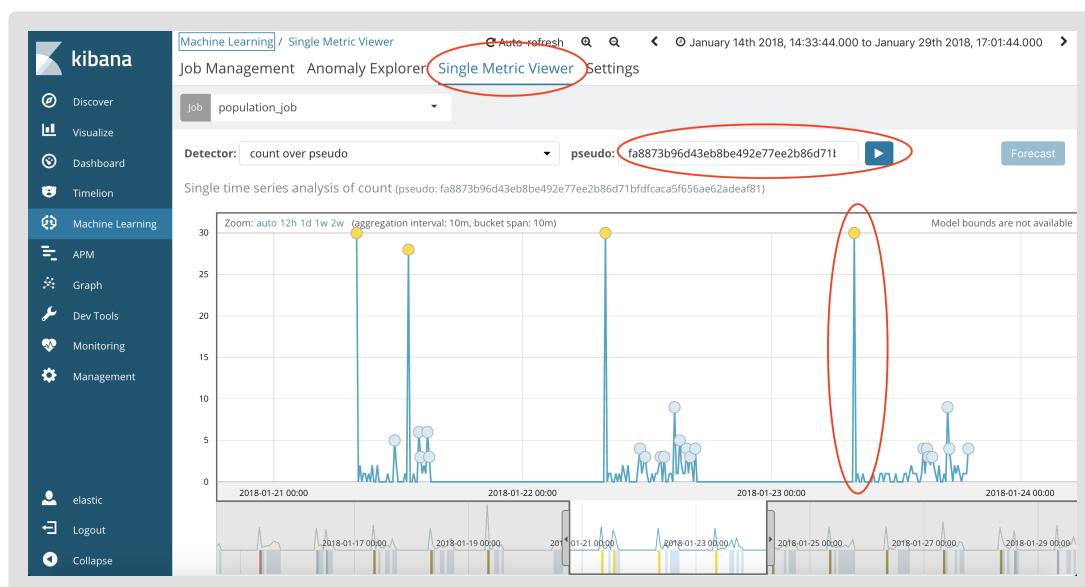
10. This job will take more time than the previous one. Once it's ready click on **View Results**. You should see a summary of the main influencers and their anomaly. Try to find the user with the following pseudo:
fa8873b96d43eb8be492e77ee2b86d71bfdccaca5f656ae62adeaf81
What is specific about this user?

Hide answer:

This user seem to have a regular pattern, and every time this user is actif, he is generating anomalies, meaning that they tweet much more than other users. image:lab4_10_1.png

11. Go into the single metric viewer and look at this user time series. When does this user publish new tweets? What's special about the way he is publishing things?

[Hide answer:](#)



When you are looking at his data, the user seem to often start tweeting around 8-9AM and stop tweeting around 5-6PM, which seem like an "office schedule". Another interesting aspect is that this user when they start tweeting seems, they tweet a lot. Maybe information that arrived during the night need to be tweeted as soon as possible, so everything is done at the first hour.

Summary: Well done!! That was the last lab! Now you know how to collect tweets, enrich those tweets and run machine learning algorithms on top of it!

End of Lab 4

© Elasticsearch BV 2015-2019. All rights reserved. Decompiling, copying, publishing and/or distribution without written consent of Elasticsearch BV is strictly prohibited.